A Survey of Communication Complexity Applications

Ryan Batubara
rbatubara@ucsd.edu

Ivy Hawks
mahawks@ucsd.edu

Darren Ho
dah103@ucsd.edu

Ciro Zhang
ciz001@ucsd.edu

Shachar Lovett
slovett@ucsd.edu

Abstract

Communication complexity studies the asymptotic communication bounds of various models of communication. The concept of communication is universal among computing: many problems in cryptography, distributed computing, game theory, and multi-agent learning, among many others, can be framed in terms of communication. In this report we survey the applications of communication complexity for a variety of fields, examining the fundamental communication problems in each and their lower and upper bounds.

Repository: https://github.com/rybplayer/DSCCapstone

1   Introduction

Communication complexity is the study of asymptotic communication bounds. For example, suppose two parties commonly named Alice and Bob each have a set of n numbers, and they wish to find if the intersection of their sets are empty. Suppose further that Alice and Bob each have access to infinite computational power. That is, we only care about the communication that happens between them.

Even with such a powerful model, communication complexity shows that such an operation must still take $\Omega(n)$ bits. Despite the power of our assumptions and the simplicity of the model, disjointness can be used to prove lower bounds for many models. For example, disjointness shows that any Turing machine with one tape must take at least $\Omega(n^2)$ time to compute disjointness written on its tape. We show this by means of a reduction.

Suppose such a machine takes $T < \Omega(n^2)$ time to compute disjointness of two sets written on a tape. Alice and Bob use it to solve the disjointness problem in less than $\Omega(n)$ communication as follows: Both Alice and Bob run the machine. Whenever the machine needs to read from Alice's part of the tape, Alice sends Bob a message to indicate what needs to be executed next. The tape can only cross between Alice and Bob's part of the tape, or vice versa, at most T/n times, because each sweep will take $\Omega(n)$ time. Therefore, using this protocol, Alice and Bob obtain a method for computing disjointess in less than $\Omega(n)$ time, a contradiction.

In this report, we survey how foundational problems in communication complexity, such as disjointness, can be used to prove bounds on a variety of problems from different fields. Four areas have become our focus: cryptography, game theory, machine learning, and agential learning. For these areas, we begin by reviewing the foundational problems in the field, before jumping into the technical proof details of bounds for such problems. We conclude by providing a cursory overview of other application areas.

2   Cryptography

2.1   Background

Cryptography is the field which deals with secret information. Specifically, it is used to hide information such that it cannot be intercepted by a third party. This originally required a symmetric key system, in which all participating parties used one secret key to decode messages, however modern systems use Public key encryption schemes, which do not require advance sharing of a secret key.

2.2   Foundational Problems

2.2.1   Public key Encryption

A Public key encryption scheme is a system which makes use of trapdoor functions, functions which can be easily computed but their inversion is enormously difficult without one key piece of information, such as multiplying two primes. The public key, or a method to compute the function, can be shared widely, and any party can pass information through it to generate an encrypted message. The reverse, however, can only be done (in efficient time) by the party which created the function, and the message can then only be read by that party.

2.2.2   Oblivious Transfer

An Oblivious transfer is a system in which one party sends some information to another party, but is not able to determine what, if any information, was sent. Thus, information can be compared in a function of some kind without the Senders ability to see what information it was, despite being in possession of it.

2.2.3   The Millionaires Problem

The Millionaire's Problem, originally proposed by Andrew Yao, is a problem for efficiently communicating under a public key cryptographic system. Suppose two Millionaires, Alice and Bob, wish to know who has more money, but do not wish to share the exact values. Yao's original solution to this problem involves the use of a garbled circuit model, in which Alice can generate a circuit which will compute the solution to the problem, and then encrypt the circuit by decomposing and anonymizing each step in the process. Alice will compute her portion of the circuit, encrypt the system such that it can only be decrypted with the possession of two inputs, allowing Bob to complete the circuit using only the portion of the truth tables which work for his input.

This solution has $\mathcal{O}(n \cdot k)$ communication complexity cost, where n is the size of the input, and k is the security parameter, which in many modern applications varies from 100 to 128.

2.3   Case Study: Garbled Circuits

Garbled Circuits is one of the most common ways to solve any Oblivious Transfer protocol, and here we will analyze it's application to the Millionaire's Problem.

2.3.1   Circuit Garbling

In the problem start, Alice and Bob each have an input $x, y \in \mathbb{R}^n$. Alice will create a circuit which computes the function (in this case, Greater Than), which will take $2 \cdot n$ inputs and

produce a Binary output. Alice will then assign two random strings of length k to each wire in the circuit, where each string matches to a value 0 or 1 for that wire. Then, using the strings which represent inputs to each wire in the circuit, can encrypt the output values for every truth tables. Each component in the circuit now has a truth table which can only be decrypted with knowledge of two wire state strings for that gate, and Alice can safely send the random strings which correspond to her inputs to Bob, who cannot decode what her inputs were merely from these strings. However, Bob must now acquire the strings which correspond to his input from Alice, without Alice knowing which inputs she has sent to Bob.

2.3.2 Oblivious Transfer Protocol

Alice will now generate a Public-Private key pair, t_{pub} and t_{priv}, as well as two random messages, m_0 and m_1, and send both m_1 and m_2 as well as t_{pub} to Bob. Bob will chose either m_1 or m_0, depending on his input, generate a random string k, and use this to then encrypt $m_b + k$, where b is the input string Bob wishes to obtain. Bob can now send this encrypted message to Alice, who then attempts to decrypt by subtracting both m_1 and m_2 from the message Bob sent, obtaining two possible options for k; k_0 or k_1. One of these values will be the original k, but Alice does not know which. She can then use these values to encrypt the random wire stings and send to Bob, who, knowing the real k, can decrypt only the one corresponding to his chosen input, thereby obtaining only the portion of the truth table which allows him to compute his input with Alice's.

Given that the circuit of the greater than function is $\Omega(n)$, and each gate will require at least $\Omega(k)$ bits, (for the encryption string length), the total cost for Alice to communicate the system to Bob is $\Omega(b \cdot k)$ bits.

2.4 Modern Applications

The garbled circuit is not only a useful model for computing two anonymous values, such as must be done in many e-commerce contexts, where a server must validate if the purchaser has sufficient funding to make the purchase, but can be adapted to compute many different functions. As the circuit complexity of this system can be modeled equivalently to any deterministic two party communication protocol, it can compute such a protocol anonymously for any valid function that two parties may wish to compute. This is mainly useful in transactional contexts, such as with blockchain, where multiple parties are not behaving with malicious intent, but nonethless wish to gain intformation some parties do not wish to disclose.

3 Game Theory

3.1 Background

The field of game theory studies the mathematical models of interactions between rational decision makers. Game theory initially addressed the interactions between two parties in a zero sum game but has grown into an umbrella term on rational decision making.

3.2 Foundational Problems

3.2.1 Algorithmic Mechanism Design

Algorithmic mechanism design (AMD) focuses on designing a mechanism where (a) each party behave in a certain way and (b) is computationally efficient. AMD differs from regular mechanism design in that mechanism design assumes infinite compute power. AMD seeks to find a mechanism that performs with certain constraints such as polynomial time.

3.2.2 Vickrey Auction

Vickrey Auction is a sealed-bid second-place auction, where the highest bidder pays the second highest price. This is a classical example of AMD because it is in the best interest of each party to bid truthfully.

Say we have three parties participating in this auction, Alice, Bob, and Charlie. Their sealed bids are:

- Alice: $50
- Bob: $100
- Charlie: $70

In this scenario, Bob wins and pays $70.

Bob cannot improve his results by changing his bid, if he raises it the price he pays does not change, and if he lowers it he risks not winning the bid. Hence in this game, there is a dominant strategy to bid truthfully.

3.2.3 Nash Equilibrium

Nash Equilibrium are a foundational problem within game theory that involve finding a state within a scenario where no player can improve their outcome by unilaterally changing their own strategy.

3.2.4   The Prisoner's Dilemma

The prisoner's dilemma is one of the simplest examples of finding a **Nash Equilibrium**.

Two suspects, Prisoner A and Prisoner B, are arrested and interrogated separately. Each prisoner has two possible actions:

- **Cooperate (C)**: remain silent
- **Defect (D)**: betray the other prisoner

The payoff matrix (with payoffs in the form (A's payoff, B's payoff)) is given by:

	B: C	B: D
A: C	$(1,1)$	$(3,0)$
A: D	$(0,3)$	$(2,2)$

The interpretation is:

- If both stay silent (C,C), each receives a light sentence: 1.
- If one defects and the other cooperates, the defector goes free (0) while the cooperator receives the harshest sentence (3).
- If both defect (D,D), they each receive a moderate sentence (2).

Even though mutual cooperation leads to a better outcome than mutual defection, the dominant strategy for each prisoner is to defect, making (D, D) the unique Nash equilibrium.

3.3   Case Study: Pure Nash Equilibrium

We will be looking to see what the communication complexity of reaching equilibrium is. More specifically, the communication model in which players initially only know their own utility functions. We want to analyze how much information must be transferred between them to jointly compute the equilibrium point. We also assume that each player follows a predetermine protocol, abstracting any other incentives of each player. We will be focusing on the communication complexity of *pure* Nash Equilibrium. A pure Nash Equilibrium being a situation where there is guaranteed to be a unique best input for a all given parties.

3.3.1   Setting for Pure Action Games

There are $n \geq 2$ players, $i = 1, 2, ..., n$. Each player i has a finite set of actions A_i with $|A_i| \geq 2$. For the analysis of pure action games, we only consider binary action games, that is, for each i, $A_i = \{0, 1\}$. Let the joint action space $A = \prod_{i=1}^{n} A_i$. Each player has a private utility function $u_i : A \rightarrow \{0, 1\}^n$ which we are assuming are finitely represented for simplicity.

For a joint action $a = (a_1, ..., a_n) \in A$ (for binary action games, a joint action is an n bit binary string), let a^{-1} denote the joint action of all players except player i. A joint action a is a **Pure Nash Equilibrium** if $u_i(a) \geq u_i(b_i, a^{-1})$ for every player i and any action $b_i \in A_i$.

That is to say that for all players choosing the alternative option would result in an equal or worse outcome for them.

3.3.2 Communication Complexity of Pure Nash Equilibrium

Theorem *Any pure Nash equilibrium procedure has communication complexity* $\Omega(2^n)$

Claim If there exists a reduction from the S-disjointness problem to an n-person pure Nash Equilibrium procedure that satisfies reducibility and constructibility properties then any pure Nash Equilibrium procedure has communication complexity of at least $|S|$ bits since it has been proven that $CC(DISJ_S) = |S|$.

We want to prove that there exists a reduction from the set disjointness problem where:

$$S_1 \cap S_2 = 0 \text{ or } S_1 \cap S_2 \neq 0$$

Our goal is to relate the disjointness problem into a game that has a pure Nash equilibrium iff the sets intersect. To do so, we use the *matching pennies reduction*.

Matching Pennies. The matching pennies game is a 2×2 zero-sum game played between two players. Each player chooses either *Heads* (H) or *Tails* (T). If the actions match, player 1 wins; if they mismatch, player 2 wins. The payoff matrix for player 1 is:

	H	T
H	1	−1
T	−1	1

and player 2 receives the negative of this payoff.

A key property of matching pennies is that it has *no pure Nash equilibrium*. This makes it useful as a "destabilizing gadget" in reductions, since attaching a matching-pennies subgame to an action profile guarantees that the profile cannot be a pure equilibrium.

For $n \geq 4$, the matching pennies reduction satisfies the reducibility and constructibility properties.

We first verify the *constructibility* property. Fix $\alpha \in \{1, 2\}$. By the definition of the matching pennies reduction, the payoff function $u_{\alpha,i}(a)$ of every player (α, i) depends only on the following information:

- whether $a \in S_\alpha$, and
- the actions $a_{h,1}$ and $a_{h,2}$ in case $a \notin S_h$

Since S_α is the private input of the agents on side α in the disjointness problem, the event $a \in S_\alpha$ is computable by those agents, and $a_{\alpha,1}, a_{\alpha,2}$ are part of the observed joint action a. Hence each payoff $u_{\alpha,i}(a)$ is computable from (a, S_α, i) alone. This establishes constructibility.

We now prove the *reducibility* property. We must show that

$$S_1 \cap S_2 \neq \emptyset \qquad \Longleftrightarrow \qquad G \text{ has a pure Nash equilibrium.}$$

7

Suppose $a \in S_1 \cap S_2$. Then by construction, every player in T_1 and T_2 receives payoff 2 at a, which is the maximal payoff any player can obtain. Thus no player can improve by deviating, and a is a pure Nash equilibrium.

Suppose $a \notin S_\alpha$ for some $\alpha \in \{1, 2\}$. Then the two designated players $(\alpha, 1)$ and $(\alpha, 2)$ play a matching pennies game at a. The matching pennies game has no pure Nash equilibrium: at every pure action profile of the two players, one of them strictly benefits from unilaterally deviating. Therefore, at a at least one of the players $(\alpha, 1)$ or $(\alpha, 2)$ has a profitable deviation, implying that a cannot be a pure Nash equilibrium of G. Hence no pure Nash equilibrium can lie outside $S_1 \cap S_2$.

Combining the two directions, we conclude that a pure Nash equilibrium exists in G if and only if $S_1 \cap S_2 \neq \emptyset$.

Since we have proven the claim, we know that the communication complexity of any pure Nash Equilibrium is $CC(DISJ_s) = |S|$. Remember that $S = \{0, 1\}^n$ so therefore $|S| = 2^n$ proving the theorem.

3.4 Broad Applications

The communication complexity of Game Theory has many varied applications and is prevalent in any situation that involves optimizing the output in any multiparty game. Game Theory draws naturally has parallels to other fields such as multi-agent and distributed systems.

4 Machine Learning

There are lots of ways to use communication complexity to prove lower bounds for data structures. Here we limit ourselves to the approximate nearest neighbor problem with static (unmodifiable and only queriable) data structures.

4.1 Background Problems

4.1.1 Nearest Neighbors (NN)

The *nearest neighbor problem* takes in a set S of n points in a metric space (X, ℓ). Most commonly this is $\mathbb{R}^d$ with the ℓ_2 norm. Here we use the *hamming cube $X = \{0, 1\}^d$* and ℓ is the Hamming distance. Let d be large, say $d = \sqrt{n}$.

The goal is to build a data structure D (as a function of $S \subseteq X$) that prepares for all possible nearest neighbor queries. A query is a poiny $q \in X$, and the end goal is to return $p^* \in X$ that minimizes $\ell(p, q)$ over all $p \in S$.

We use the Hamming cube because it is easier to map into communication complexity, though the key high-level ideas are the same for analyzing Euclidean and other metric spaces.

If computing $\ell(p, q)$ takes $O(d)$ time, then the brute force algorithm takes $O(dn)$ time. Alternatively, precomputing the answer to all possible queries takes $\Theta(d 2^d)$ space. The exact NN problem thus suffers from an exponential-size data structure in terms of d. This is known as the *curse of dimensionality*.

4.1.2 Approximate Nearest Neighbors (ANN)

The $(1 + \epsilon)$-approximate NN, where we only need to return p where $\ell(p, q) \leq (1 + \epsilon)\ell(p^*, q)$, we can do much better. Let $\epsilon = 1$ or 2, so it is not too small but still practically relevant.

The high-level idea is to hash nearby points into the same bucket, then precompute the answer for each bucket. Therefore the key challenge is to make sure nearby points hash to the same bucket.

4.2 Case Study: Decision ANN

Consider the *decision ANN* version, where given $L \in \{0, 1, \ldots, f\}$, we wish to distinguish between

1. There exists a point $p \in S$ with $\ell(p, q) \leq L$.
2. $\ell(p, q) \geq (1 + \epsilon)L$ for every $p \in S$.

If neither applies, both answers are correct.

Note that if $|S| = 1$, this is oddly similar to equality: Equality is the same as deciding between hamming distance 0 or not. The public coin protocol where Alice takes the inner product mod 2 of x with r_1 and r_2 (2 bits) and sends it to Bob always acommunication complexityepts $x = y$ and acommunication complexityepts $x \neq y$ with probability 1/4.

4.2.1 Gap-Hamming with Error

Consider the communication complexity problem where Alice and Bob want to decide if $\ell_H(x, y)$ between inputs $x, y \in \{0, 1\}^d$ is at most L or at least $(1 + \epsilon)L$. Call this ϵ-*Gap Hamming*. Define the following protocol:

1. Alice and Bob create $s = \Theta(\epsilon^{-2} \log(?))$ random strings $R = \{r_1, \ldots, r_s\} \in \{0, 1\}^d$ where $d = \theta(\epsilon^{-2})$ from public coins. Each entry is independent but not uniform: there is a $1/2L$ probability of each entry in r_i being equal to 1.
2. Alice sends s inner products of r_i with x. This is the "hash value" $h_R(x) \in \{0, 1\}^s$.
3. Bob acommunication complexityepts if any only if the Hamming distance between $h_R(y)$ and $h_R(x)$ differs by some small number of bits.

We defer the analysis of the small number of bits to the book by Roughgarden. The key idea is as follows:

1. When selecting a string r_i, we are choosing a coordinate to be relevant (be a 1, and thus contribute to the inner product later) with probabiliy $1/L$. Consider it irrelevant otherwise. Think of this as selecting relevant coordinates for this r_i hash.
2. The second stage, the inner product, is a modified equality with these coordinates only. We see that

$$\Pr_j[\langle r_j, x \rangle \not\equiv \langle r_j, y \rangle \mod 2] = \frac{1}{2}\left(\left(1 - \left(1 - \frac{1}{L}\right)^{l_H(x,y)}\right)\right).$$

This results in communication complexity $\Theta(\epsilon^{-2})$.

4.2.2   Data Structure Decision ANN

We now convert the above algorithm into a data strcuture for the $(1 + \epsilon)$ decision ANN problem.

1. Given a point set P of n points in $\{0, 1\}^d$, choose a set of R of $s = \Theta(\epsilon^{-2} \log n)$ as above.
2. Define $h_R : \{0, 1\}^d \to \{0, 1\}^s$ by setting the jth coordinate to $\langle r_j, x \rangle \mod 2$.
3. Make a table with $2^s = n^{\Theta(\epsilon^{-2})}$ buckets, indexed by s-bit strings.
4. For all $p \in P$, insert p into every bucket b which $\ell(h_R(p), b) \le (t + \frac{1}{2}h(\epsilon))s$ where $h(\epsilon) = \frac{1}{2\epsilon^2}(1 - \epsilon^{-\epsilon})$.

This data structure has probability at least $1 - \frac{1}{n}$ to give the correct answer in $n^{O(\epsilon^{-2})}$ space.

We will show later that with this level of acommunication complexityuracy this space is optimal. Note that it is possible to do better by inreasing the query time.

Note that in the real problem we do not have this value L. We would want it to be the actual NN distance of a query point q (since L counts the number of differences / coordinates checked), but this can change from one q to another. Furthermore, the data strcuture may be wrong on as many as $2^d/n$ queries. Knowing the coin flips allow an adversary to exhibit queries that will be wrong.

Fix 1: Make d copies of the data structure for each value of L^2. Answering takes $O(\log d)$ lookups. Space blows up by $\Theta(d \log d)$. Query time blows up by $\Theta(\log \log d)$.

Fix 2: Make $\Theta(d)$ copies and take the majority vote. This is wrong in at most inverse exponential of d. Here not even an adversary who knows the coin flips knows whether a particular query will be incorrect ahead of time. This blows up space and query time by $\Theta(d)$.

Fix 3: Make $\Theta(d)$ copies and choose uniformly randomly one of the copies to answer. The space blows up by $\Theta(d)$ but query time is unaffected. The probability of a correct answer is at least $1 - \Theta(\frac{1}{n})$.

Combining the three lemmas:

- Space: $O(d^2 \log d) \cdot n^{\Theta(\epsilon^{-2})}$
- Query time: $O(\epsilon^{-2} d \log n \log \log d)$.

That is, polynomial space for logarithmic query time.

4.2.3 The Cell Probe Model

We start by motivating the cell probe model. The goal is to create a database D to answer Q queries, known beforehand. The encoding scheme must work for all possible databases: In ANN, P is the database (unknown beforehand) and Q are elements of $\{0,1\}^d$ (known). Another example is set membership.

A parameter of the cell probe model is *word size* w. The design space if the number of ways to encode D as s cells of w bits each. We say s is the space used in the encoding. It must correctly answer all $q \in Q$; it does this by reading cells one at a time, giving w bits each. The query time of this algorithm is the max, over all D and all $q \in Q$, number of acommunication complexityesses to answer the queries.

w is typically large so a single element of the database can be stored in a cell, and not much larger than this.

For ANN, take w to be polynomial in $\max\{d, \log_2(n)\}$. The previous construction solves ANN cell probe with space $n^{\Theta(\epsilon^{-2})}$ and query time 1. We now show a matching lower bound in the cell-prove model; thus constant query time can only be achieved by encodings that us $n^{\Omega(\epsilon^{-2})}$ space.

4.2.4 Query Database

Consider the *query database* problem. Alice has a query q and Bob gets database D. We wish to compute the answer to q on the database D.

Let there be a cell-probe encoding of query database with word size w, space s, and query time t. Then there is a communication complexity protocol with communication complexity at most

$$\underbrace{t \log_2(s)}_{Alice} + \underbrace{tw}_{Bob}$$

Alice simulates the query-answering algorithm, sending $\log_2(s)$ to Bob for every cell requested. Bob sends back w bits to describe the contents of the selected cell. They need to go back and forth only t times (since the query time is also the number of $\Theta(1)$ queries, i.e. t queries).

[*Richness Lemma*] Let $f : X \times Y \to \{0,1\}$ be a Boolean function with corresponding $X \times Y$ boolean matrix $M(f)$. Assume that

1. $M(f)$ has at least v columns that each have at least u 1-inputs.
2. There is a deterministic protocol that computes f in which Alice and Bob send a and b bits respectively.

Then $M(f)$ has a 1-rectangle $A \times B$ with $|A| \geq u/2^a$ and $|B| \geq u/2^{a+b}$.

For every $\epsilon, \delta > 0$ and sufficiently large n, in every communication complexity protocol that solves $(\frac{1}{\epsilon^2}, n)$-disjointness with a universe of size $2n$, either

1. Alice sends at least $\delta/\epsilon^2 \log_2 n$ bits, or
2. Bob sends at least $n^{1-2\delta}$ bits.

$(\frac{1}{\epsilon^2}, n)$-disjointness reduces to the query database problem for the decision problem of ANN with $(1 + \epsilon)$.

Every data structure for the decision version of $(1 + \epsilon)$-ANN with query time $t = \Theta(1)$ and word size $w = O(n^{1-\delta})$ for $\delta > 0$ uses space $s = n^{\Omega(\epsilon^{-2})}$.

Using the cell probe model, it is possible to convert our lower bounds for DANN into lower bounds for ANN. However, a full proof is beyond the scope of this text, and readers are recommended to see Roughgarden's book and the bibliography for more details.

5 Agential Learning

5.1 Background: Multi-Agent Reinforcement Learning

Multi-agent reinforcement learning (MARL) studies environments where several learning agents act simultaneously. Each agent i receives a private observation o_i, takes an action a_i, receives a reward, and updates its own policy:

$$a_i \sim \pi_i(o_i).$$

Policies are typically represented as neural networks.

5.1.1 Motivation

In many real-world systems, such as fleets of autonomous cars, distributed robots, or cooperative drones, agents cannot observe the full global state. Limited sensing creates uncertainty, and achieving coordinated behavior requires exchanging information.

However, communication is expensive:

- Bandwidth is limited.
- Communication introduces latency.
- Regulatory or safety constraints may restrict what agents can share.

Thus, communication is not free. We must understand:

What is the minimum number of bits agents must exchange to learn effectively?

Communication complexity offers a rigorous framework for proving *lower bounds* on required communication.

5.1.2   Example Environment

Agents explore a grid. Their objective is to collectively visit all squares.

Rewards:

- +1: discover a new square
- −0.5: illegal action
- −0.5: revisiting an old square
- −1: no movement
- +200: full coverage of the grid

Because each agent only observes nearby cells, communication is needed to avoid redundancy (e.g., multiple agents going to the same place).

5.1.3   Cooperative MARL: Learning with Shared Parameters

We consider M agents solving the same task and sharing a global model. At each episode:

1. Each agent computes a local score or gradient estimate.
2. Updates to the shared model occur only if the new estimate is significantly different.

5.1.4   Communication Complexity Result

Let d = dimension of model parameters, H = episode horizon, and M = number of agents. Then total communication per update can be bounded by:

$$O(dHM),$$

or, in settings where aggregation happens per step instead of per episode:

$$O(dM).$$

These bounds show that as the number of agents increases, communication cost scales linearly unless compression or sparsification is introduced.

5.2 Case Study: Learning from Distributed Datasets

We now turn to a different setting: *distributed supervised learning*. Two parties hold different datasets:

- - Alice has SA,
- - Bob has SB.

Their goal: Find a classifier H consistent with both SA and SB, while communicating as little as possible. This models distributed systems, federated learning, and privacy-preserving ML.

5.2.1 The Realizability Question

First question: Does a consistent classifier exist at all?

In many hypothesis classes (e.g., linear separators in $\mathbb{R}^d$), determining realizability requires communicating on the order of:

$$O(d)$$

bits. This is the *minimum* information needed to ensure that no contradiction exists between S_A and S_B.

5.2.2 Learning the Classifier

Once realizability is established, the next step is to learn a classifier. A typical protocol could be:

1. Alice sends a *weak hypothesis h* (low-accuracy model trained on S_A).
2. Bob checks h on his data S_B:
   - If h performs well, training is complete.
   - Otherwise Bob sends examples that h misclassifies.
3. Alice improves the classifier using Bob's counterexamples.
4. Iterate until both sides are satisfied.

This resembles boosting, where each interaction reduces error. The total communication requirement is

$$O\left(d\log\left(\frac{1}{\epsilon}\right)\right),$$

where ϵ is the target classification error.

5.2.3 Interpretation

This provides a fundamental lower bound: No matter what distributed learning algorithm you design, at least $\Omega(d)$ bits must be exchanged to jointly learn a d-dimensional classifier.

You cannot cheat this by clever optimization—communication is inherently tied to dimensionality. This contrasts with MARL, where communication lower bounds depend on:

- number of agents M
- episode horizon H
- parameter dimension d
- precision required for coordinated behavior

6 Other Areas

In order to drive home the flexibility of communication complexity for studying algorithmic bounds, we list a few additional application fields for communication complexity below.

6.1 Circuit Design

Boolean circuits are a natural model for computation. Oftentimes, we may be interested in the number of gates, or depth of a circuit, as measures of its complexity. A *formula* is a circuit whose underlying decision graph on inputs is a tree. That is, every circuit of depth d can be turned into a formula of depth d.

Communication complexity relates to the study of circuits via Krachmer-Wigderson games: In the game, players have a function $f : \{0, 1\}^n \to \{0, 1\}$ such that Alice gets $x \in f^{-1}(0)$ and Bob gets $y \in f^{-1}(1)$. The goal is to find $i \in [n]$ such that $x_i \neq y_i$. As it turns out, a "circuit of depth d computing f yields a length d deterministic protocol for the associated game. If the circuit is monotone, the protocol solves the monotone game" (Rao and Yehudayoff, Lemma 9.2). In other words, circuit depth is equivalent to communication complexity, because both are measures for the depth of a protocol tree.

Furthermore, if f is solvable with d bits of communication, then there is a circuit of depth d computing f (Rao and Yehudayoff, Lemma 9.3). Therefore, most surprisingly, the study of how a *single* circuit can solve problems can be modeled by the communication of two parties, each with infinite compute power!

6.2 Data Structures

Communication complexity can also be used to show bounds for the analysis of data structures. For example, consider the problem of predecessor search. That is, maintaining a set of $S \subseteq [n]$ with deletion, addition, and predecessor all being fast. It is well known that binary search trees can do addition and deletion in $O(\log n)$ time, and queries in $O(\log \log n)$ time.

The van Emde Boas tree allows us to improve on this, with $O(\log \log n)$ updates. Using communication complexity, we can prove that this data structure is optimal. This can be done by framing the sorting and prefix sum problems as communication problems, and,

similar to the proof of the turing machine reduced into disjoitness in the introduction, we can show van Emde Boas is tight optimal. See Rao and Yehudayoff, theorems 11.3 and 11.4 for more details.

7 Conclusion

In conclusion, communication complexity provides a flexible framework to prove bounds, construct algorithms, and analyze problems in a range of applicable fields. Though theoretical in nature, assuming infinite compute power from its players, the model is suprisingly applicable, with plenty of optimality proofs, algorithms, and insights that would be difficult to prove using other tools.

References

Even, Shimon, Oded Goldreich, and Abraham Lempel. 1985. "A randomized protocol for signing contracts." *Commun. ACM* 28 (6), p. 637–647. [Link]

Hart, Sergiu, and Yishay Mansour. 2010. "How long to equilibrium? The communication complexity of uncoupled equilibrium procedures." *Games and Economic Behavior* 69 (1): 107–126. [Link]

Kohl]1186828 **Kohl, Lisa Maria.** [2019]. *On improving Communication complexity in cryptography /*. Karlsruhe, [Link]

Rao, Anup, and Amir Yehudayoff. 2015. *Communication Complexity*. 11 of Foundations and Trends in Theoretical Computer Science Now Publishers. [Link]

Roughgarden, Tim. 2015. "CS261: Optimization and Algorithmic Complexity — Lecture Notes." [Link]

Yao, Andrew Chi-Chih. 1986. "How to generate and exchange secrets." In *27th Annual Symposium on Foundations of Computer Science (sfcs 1986)*. [Link]